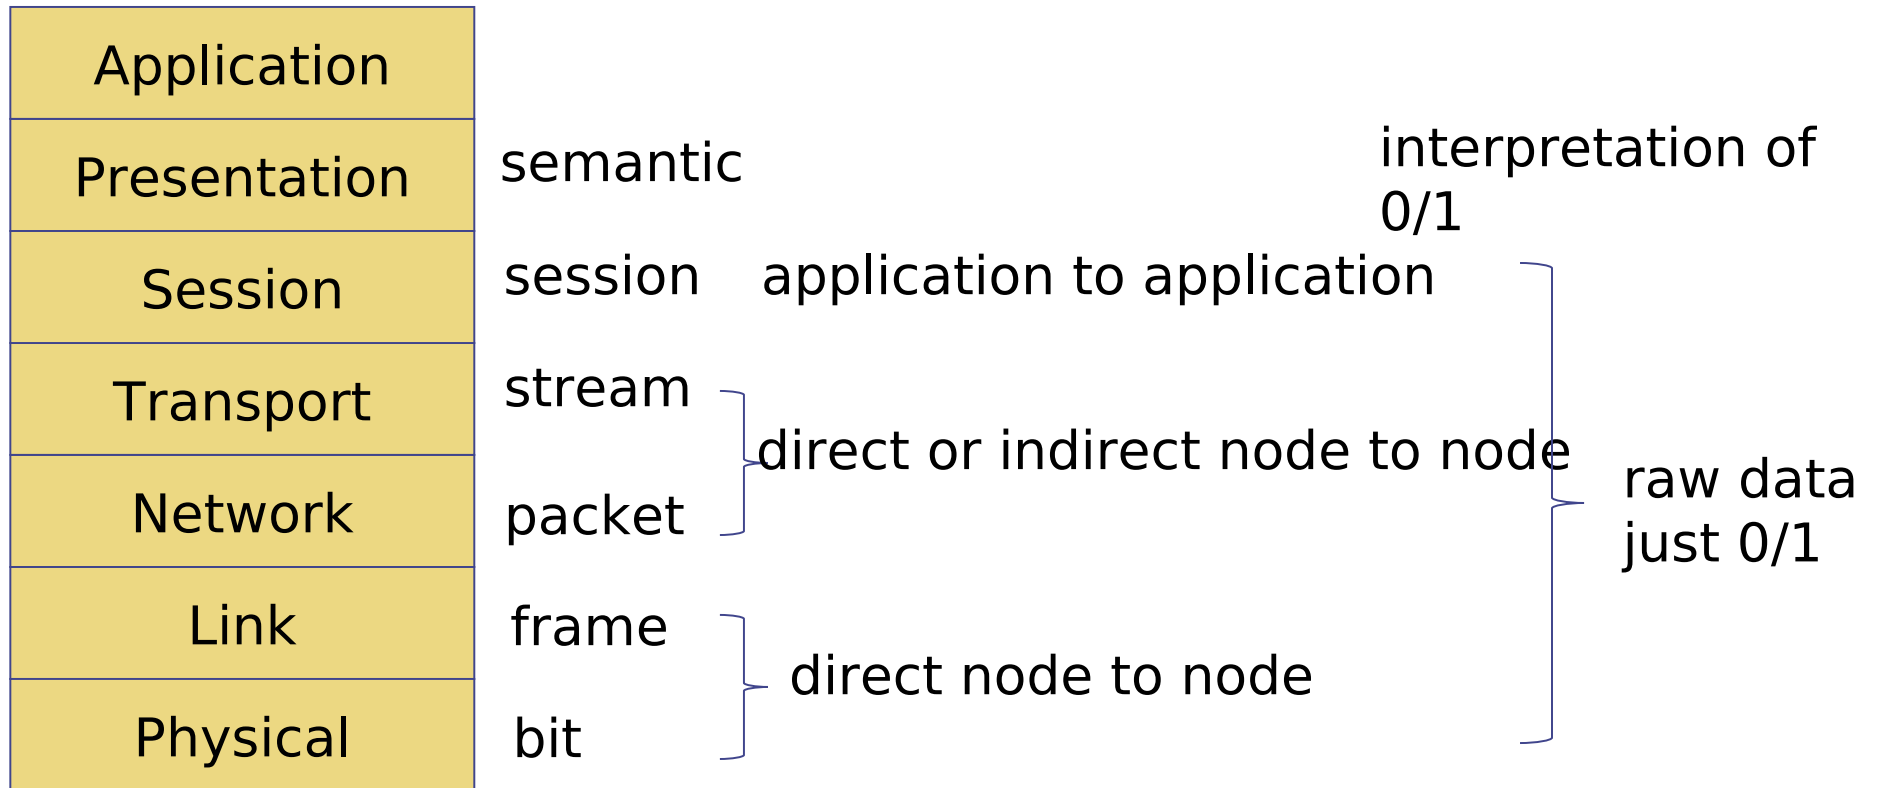
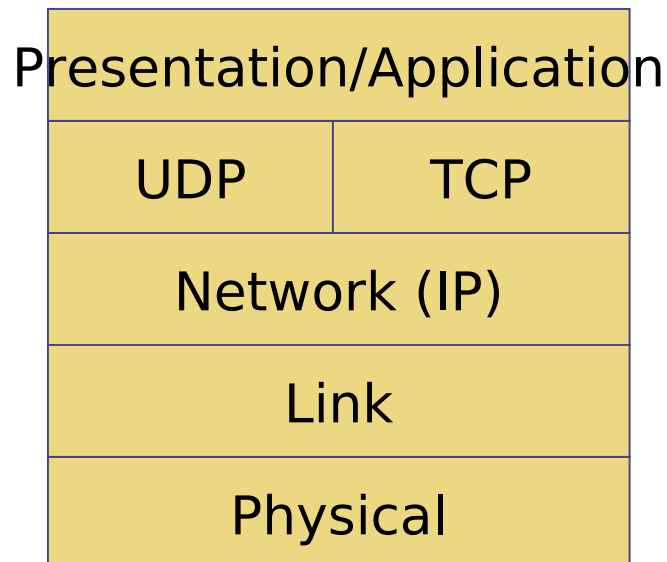
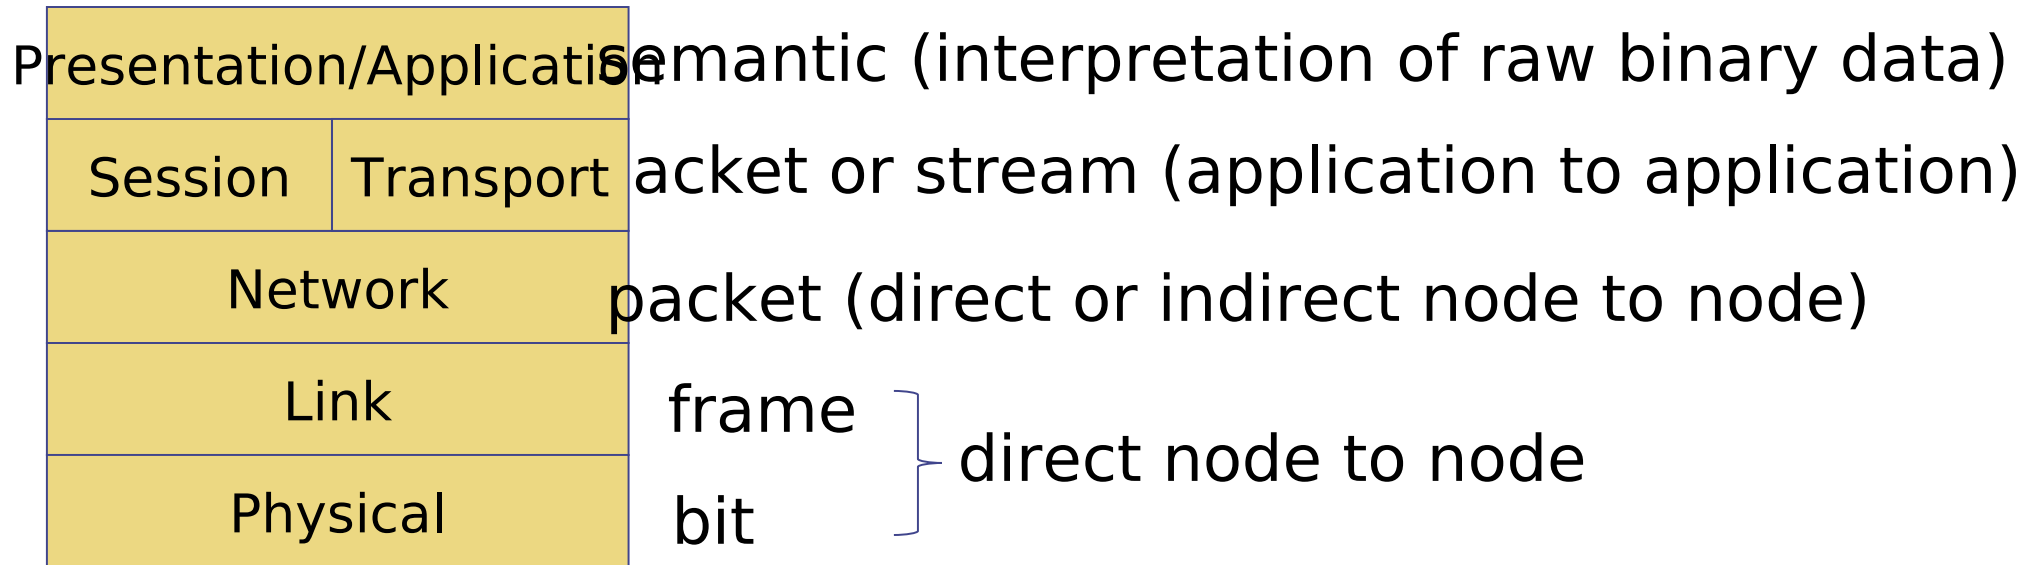


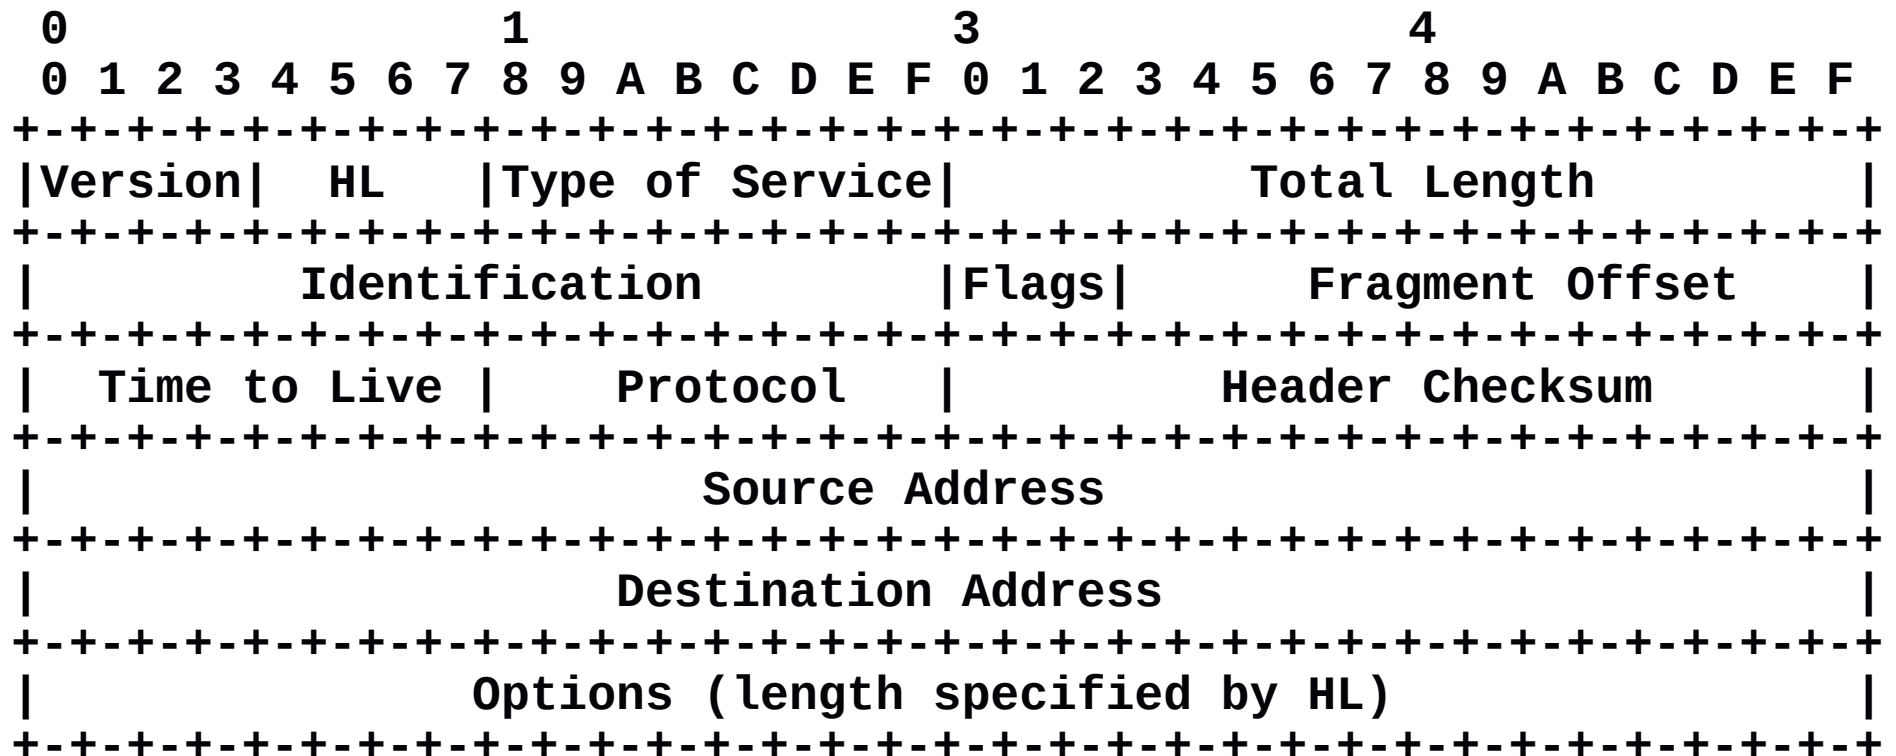
OSI 7-Layer Model



OSI 5-Layer Model



IPv4 Header



Version: 4

HL: header length, minimum 5 (x32 bits)

Type of Service: priority (real-time voice etc.), but now may be redefined

ID + Fragment offset: the group ID of a datagram, which may be fragmented in transmission

Flags: Specify if fragmentation is allowed

Protocol: the protocol used in the data portion, helping interpret the data portion (e.g., udp or tcp headers)

Internet Protocol (IP)

- Service
 - Send/receive a packet to/from a remote machine
- Interface 1: **IP_Send(dest, buf)**
 - Create a packet (**IP header + buf**)
 - Find a routing path to host **dest**
 - Send the data in **buf** to host **dest**
- Interface 2: **IP_Recv(buf)**
 - Receive an IP packet
 - Deposit the packet into **buf**
 - Return the packet size

Problems of IP

- The interface is called by all applications in the same machine
 - How to decide which application gets which packets?
- IP Packets have limited size
 - Each packet can be no more than 64K bytes
- IP is connectionless and does not guarantee packet delivery
 - Packets can be delayed, dropped, reordered, duplicated
- No congestion control

Implementation of UDP and TCP

CS587x Lecture 2
Department of Computer Science
Iowa State University

User Datagram Protocol (UDP)

- Goal
 - Provide application-to-application communication
- Idea: Concept of port
 - Each connection links to a specific *port*
 - ♦ (srcIP, srcPort, dstIP, dstPort) uniquely identifies each connection
- Interface
 - `UDP_Send(dstIP, buf, port)`
 - `UDP_Recv(buf, port)`
- Service
 - Send datagram from (srcIP, srcPort) to (dstIP, dstPort)
 - Service is unreliable, but error detection possible

UDP Send



- **UDP_Send(dstIP, buf, port)**
 1. **Prepare header**
 - 1) **source port (usually any available port)**
 - 2) **destination port**
 - 3) **UDP length**
 - 4) **UDP checksum**
 2. **IP_Send(dstIP, header+buf)**


0	15	16	31
Source Port		Destination Port	
UDP length		UDP checksum	
Payload (variable)			

Header: 8 bytes

UDP Receive



Port List



■ **UDP_Recv(buf, port)**

1. **IP_recv(buf)**
2. Get **port** information from the udp header encoded in **buf**
3. Any listener on **port**?
 - a) Yes, drop the payload to the message queue of the listener and wake it up (if it is waiting)
 - b) No, discard the packet

0	15	16	31
Source Port		Destination Port	
UDP length		UDP checksum	
Payload (variable)			

Problems of Port

- Each connection links to a specific *port*
 - (srcIP, srcPort, dstIP, dstPort) uniquely identifies each connection
- Totally there are 65535 ports
 - *Well known ports* (0-1023): everyone agrees which services run on these ports
 - ◆ e.g., ssh:22, http:80, snmp: 24
 - ◆ Access to these ports needs administrator privilege
 - *Ephemeral ports* (most 1024-65535): given to clients
 - ◆ e.g. chatclient gets one of these
 - ◆ Port contention rises

Transmission Control Protocol (TCP)

- Design Goals
 - Provide application-to-application communication
 - Messages can be of arbitrary length
 - Provide reliable and in-order delivery
 - Provide congestion control and avoidance
- Basic idea to address reliability
 - Using timer: sending a packet and waiting for an acknowledgement from the receiver
 - ◆ if the ACK is received within a given time period, the packet is received;
 - ◆ otherwise, the packet is assumed to be lost

TCP Header

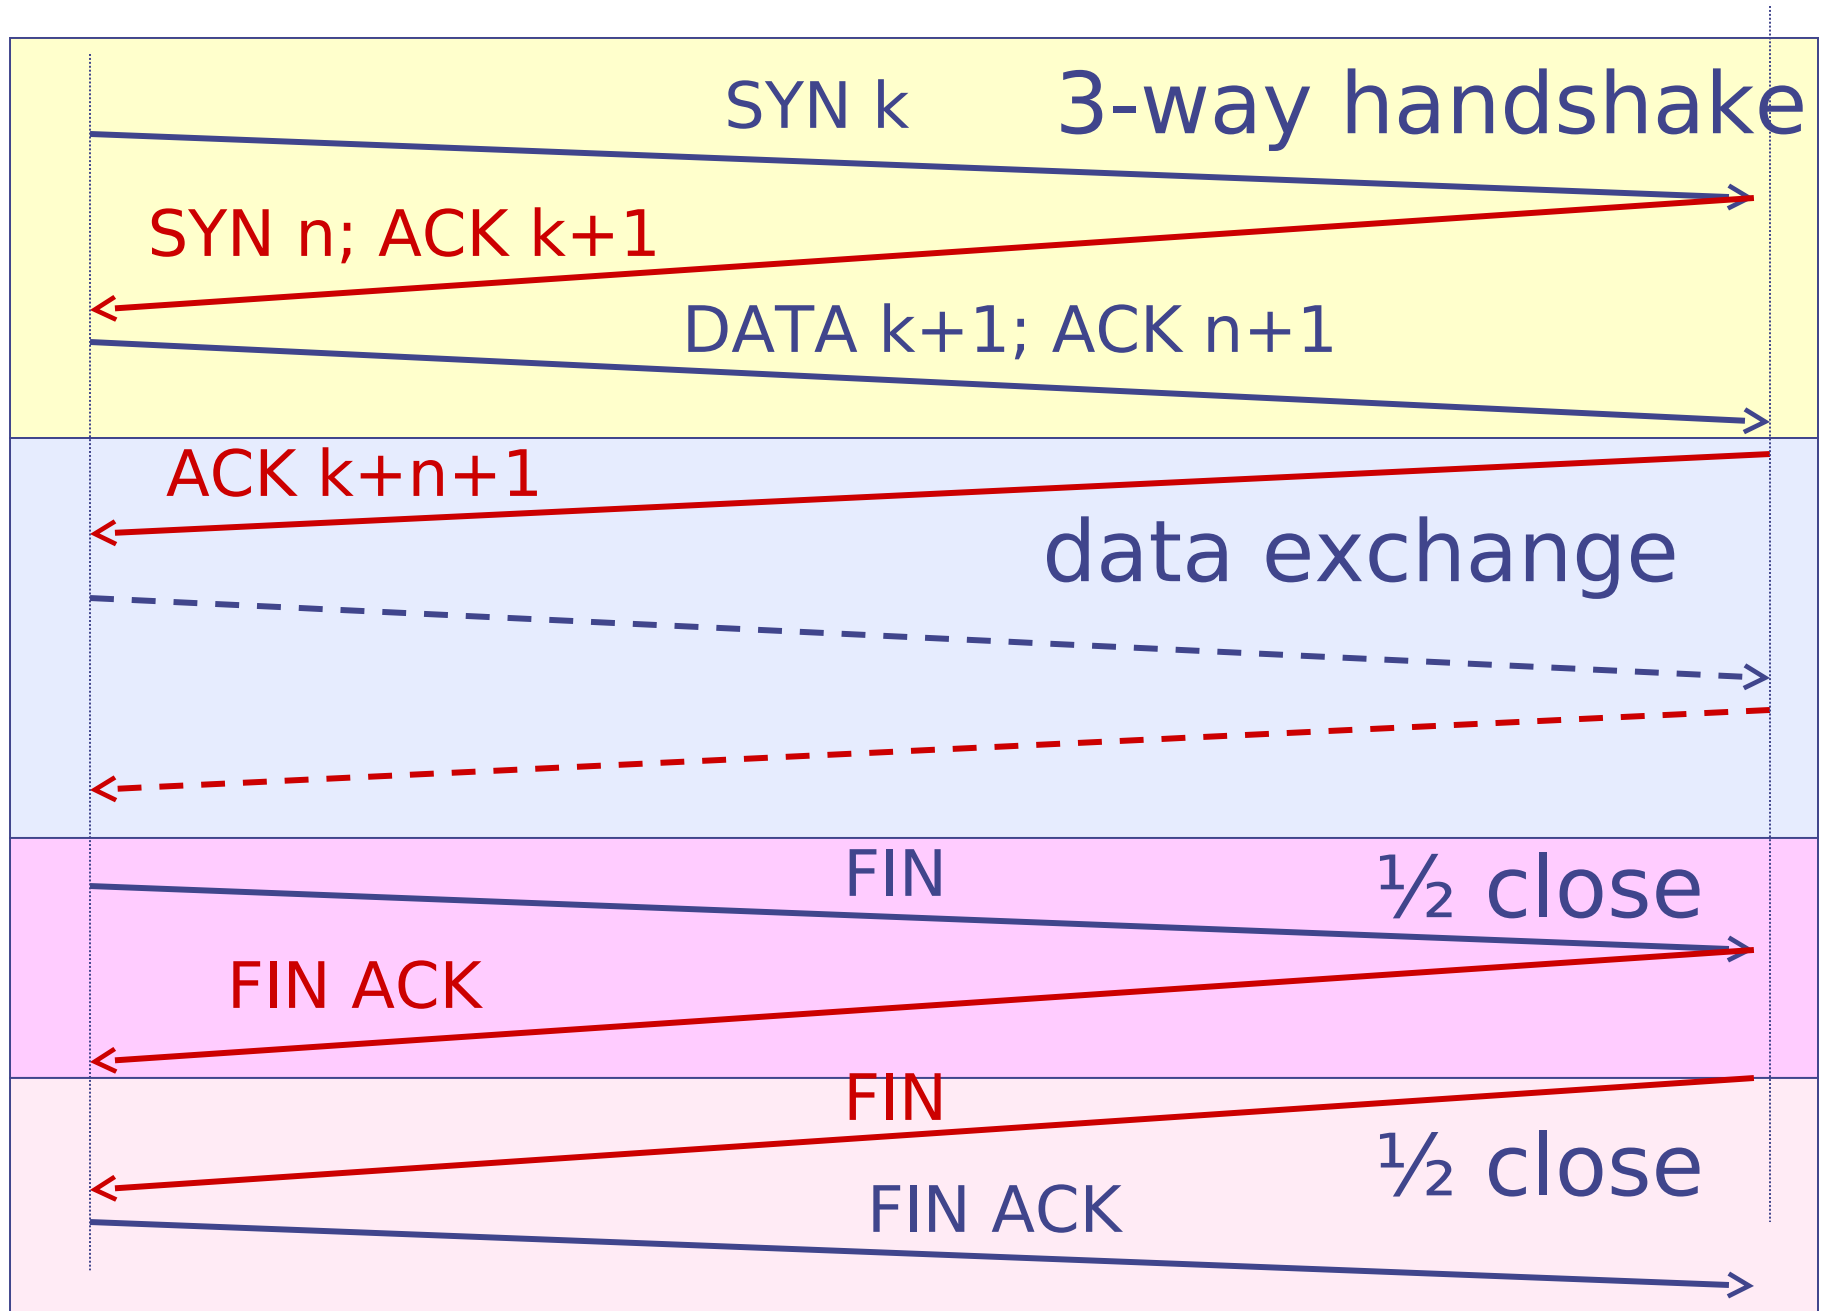
0	4	10	16	31
Source port			Destination port	
Sequence number				
Acknowledgement				
HdrLen		Flags	Advertised window	
Checksum			Urgent pointer	
Options (variable)				
Payload (variable)				

- Sequence number, acknowledgement, and advertised window – used by *sliding-window based flow control*
- Flags:
 - SYN – establishing a TCP connection
 - FIN – terminating a TCP connection
 - ACK – set when Acknowledgement field is valid
 - URG – urgent data; Urgent Pointer says where non-urgent data starts (not defined by standard, but specific to implementation)
 - PUSH – don't wait to fill segment
 - RESET – abort connection

TCP Implementation

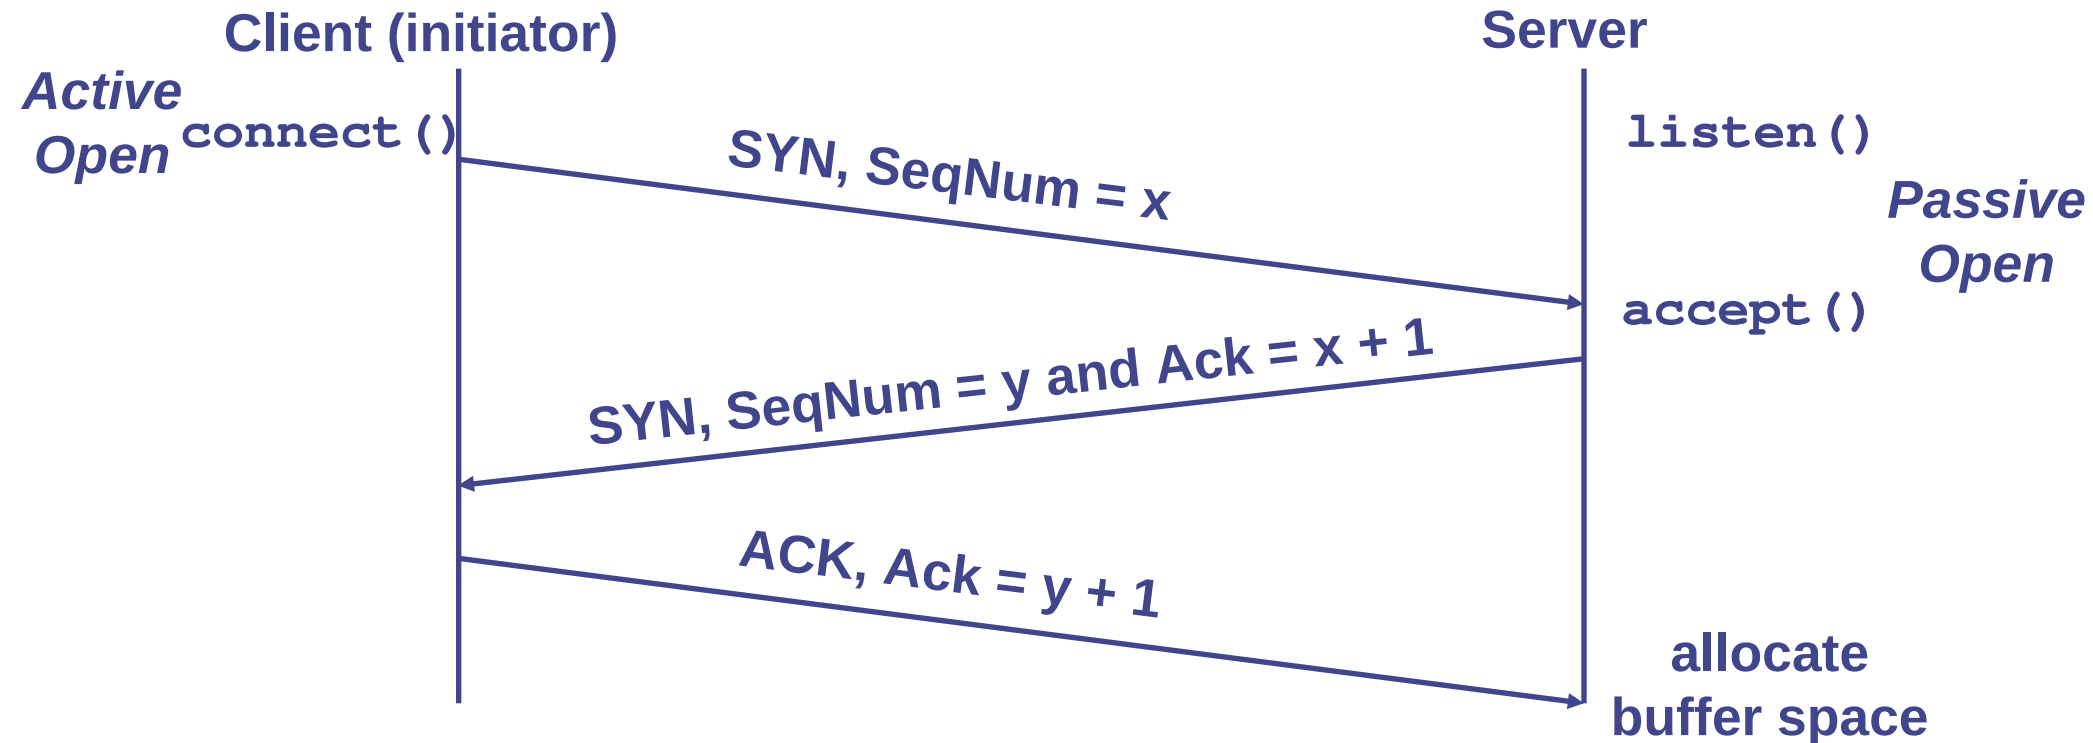
- Start a connection
- Reliable byte stream delivery from (srcIP, srcPort) to (dstIP, dstPort)
 - Indication if connection fails: Reset
- Terminate the connection

Connect/Exchange/ Terminate



Connection: 3-Way Handshake

- Three messages (i.e., syn, syn-ack, ack) are exchanged before data transmission
- Exchange sequence number, total buffer size and the size of the largest segment that can be handled at each side



The initial SeqNums must be randomized!!

Why 3WH?

- Three-way handshake adds 1 RTT delay
 - Expensive for small connections such as RPC
- Why?
 - Congestion control: SYN (40 byte) acts as cheap probe
 - ◆ Both sender and receiver has now had one round of trip, which tells some information such as their distance.
 - Protects against delayed packets from other connection (would confuse receiver)

DoS: TCP SYN Flooding

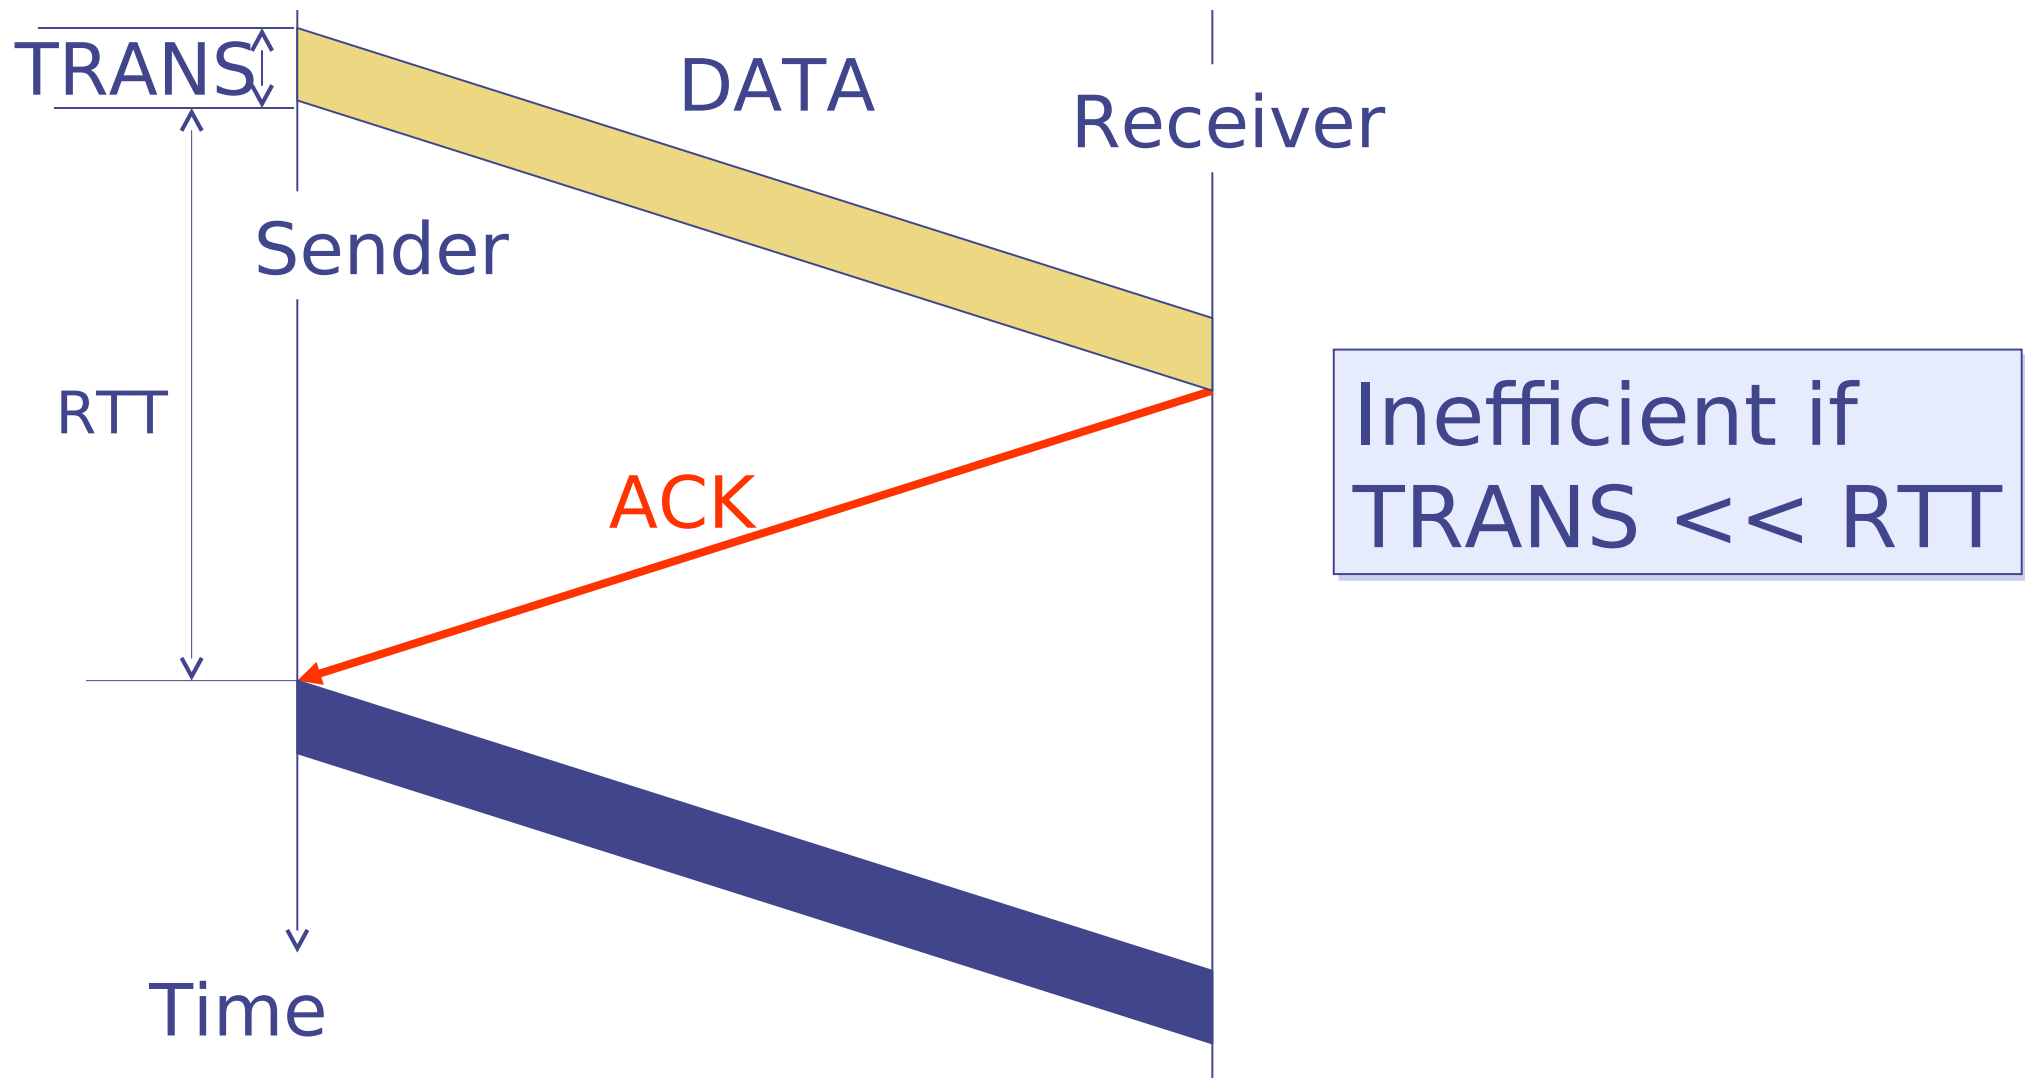
- How it works: exhausting system resources
 - Using a faked IP address
 - Initiates a TCP connection to a server with a faked IP address
 - ◆ Sends a SYN message
 - ◆ The server responds with a SYN-ACK
 - ◆ Since the address does not exist, the server needs to wait until time out
 - The server never receives the ACK (the final stage of the TCP connection)
 - Repeat with a new faked IP address
 - ◆ Repeat at a pace faster than the TCP timeouts release the resources
 - ◆ All resources will be in use and no more incoming connection requests can be accepted.
- Some common ways to prevent
 - Install firewall
 - ◆ choose deny instead of reject, which sends a message back to the sender
 - Close all ports that are not in use
 - Deny requests from unusual IP addresses
 - ◆ Private address
 - ◆ Multicast address, etc.

Termination: 4-Way Handshake

- Four messages (FIN, ACK, FIN, ACK) are exchanged to terminate a connection
 - FIN from B to A
 - ◆ B does not transmit any new data, but is still responsible for any corrupted data
 - ACK from A to B
 - FIN from A to B
 - ◆ After reading all of the bytes from B, A sends FIN to B
 - ACK from B to A
 - ◆ The connection is formally closed

Exchange: Stop & Wait


- Send; wait for ack
- If timeout, retransmit; else repeat



Exchange: Go-Back-n (GBN)

- Sliding Window Protocol
 - Transmit up to n unacknowledged packets/bytes
 - If timeout for $ACK(k)$, retransmit $k, k+1, \dots$

Example without errors




$n = 9$ packets in
one RTT instead of
1

→ Fully efficient

Example with errors

Window size = 3 packets



Observations

- Pros:
 - It is possible to fully utilize a link, provided the sliding window size is large enough.
Throughput is $\sim (w/RTT)$
 - Stop & Wait is like $w = 1$.
- Cons:
 - Sender has to buffer all unacknowledged packets, because they may require retransmission
 - Receiver may be able to accept out-of-order packets, but only up to its buffer limits


Sliding Window Size

- What size should the window be?
 - Too small:
 - ◆ Inefficient, degenerated to S&W when $w=1$
 - Too large:
 - ◆ more buffer required for both sender and receiver
 - ◆ Transmitting too fast results in network congestion and packet lost
- Congestion control
 - Slow-start phase
 - ◆ Initially set to be 1 or 2
 - ◆ Increase the window by 1 for each ACK received (this results in multiplicatively increase)
 - Congestion-avoidance phase
 - ◆ The window is increased by only 1 at a time after it is larger than the slow-start threshold (i.e., half of the size that causes congestion)
 - In the case a packet is lost, the window is decreased by half ($\text{window} / 2$).


Timer

- The sender needs to set timers in order to know when to retransmit a packet that may have been lost
- How long to set the timer for?
 - **Too short**: may retransmit before data or ACK has arrived, creating duplicates
 - **Too long**: if a packet is lost, will take a long time to recover (inefficient)

Illustrations



Timer too long



Timer too short

Adaptation

- The amount of time the sender should wait is about the round-trip time (RTT) between the sender and receiver
 - For link-layer networks (LANs), this value is essentially known
 - For multi-hop WANS, rarely known
- Must work in both environments, so protocol should adapt to the path behavior
- Measure successive ack delays $T(n)$
Set timeout = average + 4 deviations

Questions of ACKs

- **What** exactly should the receiver ACK?
- Some possibilities:
 - ACK **every packet**, giving its sequence number
 - use ***cumulative ACK***, where an ACK for number n implies ACKS for all $k < n$
 - use ***negative ACKs*** (NACKs), indicating which packet did not arrive
 - use ***selective ACKs*** (SACKs), indicating those that did arrive, even if not in order

Parallel/Concurrent FTP?


- Multi-Source Downloading
 - A large file may be available on multiple servers
 - ◆ The connections to these servers may not be reliable
 - To speed up downloading, a client may download the file from several servers
 - ◆ Subject to the limitation of client download bandwidth
 - What factors to consider?
 - ◆ Which part of the file should be downloaded from a server
 - ◆ What happens if some server is down?
 - ◆ How about disk I/O cost?

Summary

- UDP: application-to-application
- TCP: Reliable Byte Stream
 - Connect (3WH); Exchange; Close (4WH)
 - Reliable transmissions: ACKs...
 - S&W not efficient → Go-Back-n
 - What to ACK? (cumulative, ...)
 - Timer Value: based on measured RTT

Review of TCP/IP Suite

- IP header → used for IP routing, fragmentation, error detection, etc.
- UDP header → used for multiplexing/demultiplexing, error detection
- TCP header → used for multiplexing/demultiplexing, data streaming, flow and congestion control



Reading Materials

1. https://en.wikipedia.org/wiki/Transmission_Control_Protocol
2. https://simple.wikipedia.org/wiki/User_Datagram_Protocol

Review Questions

1. How can the unreliable IP be leveraged to implement reliable TCP?
2. Explain the three stages of TCP transmission: handshake, data exchange, and termination.
3. In handshake, it is crucial to set the sequential number to be a random number. Why? If the number is set to be a constant, say 0, what could go wrong?
4. There are two ways to implement data exchange: Stop and Wait, and Go-back-n. Explain their pros and cons.
5. Explain the concept of sliding windows. What make(s) it difficult to set an appropriate sliding window?
6. Explain one kind of deny-of-service (DOS) attack on TCP.